

Edge Intelligence

P Sushma

25MML0050

Task 1:

- Analysing an image dataset MNIST, and applying basic processing operations.
- Applied Convolution neural network on the following dataset.
- Saving the model using pickle model.

```
[2]: import tensorflow as tf
from tensorflow.keras import layers, models
import pickle

(train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()

x_train, y_train = train_images[:600], train_labels[:600]
x_test, y_test = test_images[400:], test_labels[400:]
x_train = x_train.reshape((600, 28, 28, 1)).astype('float32') / 255.0
x_test = x_test.reshape((400, 28, 28, 1)).astype('float32') / 255.0

[3]: model = models.Sequential([
    layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    layers.MaxPooling2D((2, 2)),

    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),

    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(10, activation='softmax')
])

C:\Users\sushm\anaconda3\envs\ml_lab\lib\site-packages\keras\src\layers\convolutional\base_conv.py:113: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.
super().__init__(activity_regularizer=activity_regularizer, **kwargs)

[4]: model.compile(optimizer='adam',
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

print("Training CNN on 600 images...")
model.fit(x_train, y_train, epochs=10, batch_size=32)

Training CNN on 600 images...
Epoch 1/10
19/19 3s 17ms/step - accuracy: 0.2834 - loss: 2.1807
Epoch 2/10
19/19 0s 11ms/step - accuracy: 0.6867 - loss: 1.2433
Epoch 3/10
19/19 0s 14ms/step - accuracy: 0.7901 - loss: 0.6412
Epoch 4/10
19/19 0s 20ms/step - accuracy: 0.8976 - loss: 0.3552
Epoch 5/10
19/19 0s 15ms/step - accuracy: 0.9333 - loss: 0.2504
Epoch 6/10
19/19 0s 12ms/step - accuracy: 0.9436 - loss: 0.1910
Epoch 7/10
19/19 0s 15ms/step - accuracy: 0.9583 - loss: 0.1598
Epoch 8/10
19/19 0s 14ms/step - accuracy: 0.9804 - loss: 0.0949
Epoch 9/10
19/19 0s 10ms/step - accuracy: 0.9836 - loss: 0.0825
Epoch 10/10
19/19 0s 13ms/step - accuracy: 0.9826 - loss: 0.0591
[4]: <keras.src.callbacks.history.History at 0x2b90c2adac0>

[5]: print("\nEvaluating CNN on 400 images...")
test_loss, test_acc = model.evaluate(x_test, y_test, verbose=1)
print(f"CNN Test Accuracy: {test_acc * 100:.2f}%")

Evaluating CNN on 400 images...
13/13 0s 12ms/step - accuracy: 0.9174 - loss: 0.2569
CNN Test Accuracy: 90.50%
```

```
[7]: with open('mnist_cnn_model.pkl', 'wb') as f:
    pickle.dump(model, f)
print("\nCNN model saved to mnist_cnn_model.pkl")

CNN model saved to mnist_cnn_model.pkl

[8]: import pickle
import os
import tensorflow as tf
import numpy as np
with open('mnist_cnn_model.pkl', 'rb') as f:
    loaded_model = pickle.load(f)

print("--- Model Loaded Successfully ---")
--- Model Loaded Successfully ---

[9]: first_layer = loaded_model.layers[0]
weights, biases = first_layer.get_weights()

print(f"Layer Name: {first_layer.name}")
print(f"Weights Shape (Filter H, Filter W, Channels, Num Filters): {weights.shape}")
print(f"Biases Shape: {biases.shape}")
print("\nSample Weights (first 3x3 filter kernel):\n", weights[:, :, 0, 0])

Layer Name: conv2d
Weights Shape (Filter H, Filter W, Channels, Num Filters): (3, 3, 1, 32)
Biases Shape: (32,)

Sample Weights (first 3x3 filter kernel):
[[-0.15950842 -0.04815693 -0.06375122]
 [-0.08455123 -0.01896419 -0.0415653 ]
 [ 0.17882669 -0.01614539  0.18866213]]]

[11]: (train_images, train_labels), (test_images, test_labels) = tf.keras.datasets.mnist.load_data()
x_train = train_images[:600].reshape((600, 28, 28, 1)).astype('float32') / 255.0
y_train = train_labels[:600]
x_test = test_images[:400].reshape((400, 28, 28, 1)).astype('float32') / 255.0
y_test = test_labels[:400]

[12]: print("\n--- Resuming Training for 2 More Epochs ---")
loaded_model.fit(x_train, y_train, epochs=2, batch_size=32)

print("\n--- Running Evaluation ---")
test_loss, test_acc = loaded_model.evaluate(x_test, y_test, verbose=1)
print(f'Test Accuracy: {test_acc * 100:.2f}%')

file_size = os.path.getsize('mnist_cnn_model.pkl') / (1024 * 1024) # Convert to MB
print(f'\nModel File Size on Disk: {file_size:.2f} MB')

--- Resuming Training for 2 More Epochs ---
Epoch 1/2
19/19 ━━━━━━━━ 2s 20ms/step - accuracy: 0.9830 - loss: 0.0552
Epoch 2/2
19/19 ━━━━━━━━ 0s 16ms/step - accuracy: 0.9992 - loss: 0.0223

--- Running Evaluation ---
13/13 ━━━━━━━━ 1s 12ms/step - accuracy: 0.9589 - loss: 0.1439
Test Accuracy: 93.50%

Model File Size on Disk: 1.43 MB
```

```
[13]: model.summary()

Model: "sequential"



| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D)                | (None, 26, 26, 32) | 320     |
| max_pooling2d (MaxPooling2D)   | (None, 13, 13, 32) | 0       |
| conv2d_1 (Conv2D)              | (None, 11, 11, 64) | 18,496  |
| max_pooling2d_1 (MaxPooling2D) | (None, 5, 5, 64)   | 0       |
| flatten (Flatten)              | (None, 1600)       | 0       |
| dense (Dense)                  | (None, 64)         | 102,464 |
| dense_1 (Dense)                | (None, 10)         | 650     |



Total params: 365,792 (1.40 MB)

Trainable params: 121,930 (476.29 KB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 243,862 (952.59 KB)
```

Task 2:

The following steps for gathering and organizing data using the Edge Impulse platform:

1. Account Setup

Begin by visiting the Edge Impulse website to create a new user account. Once registered and logged into the Edge Impulse Studio, you will have access to the dashboard where you can manage projects, datasets, and training configurations.

2. Accessing Data Acquisition

Navigate to the Data Acquisition tab within your project dashboard. This section serves as the central hub for collecting and uploading the raw inputs—such as images, audio, or sensor readings—required to train your model.

3. Connecting a Device

Click on the "Connect Data" option. This feature allows you to link external hardware (such as a development board, computer, or smartphone) directly to the platform for real-time data streaming.

4. Smartphone Integration via QR Code

Edge Impulse will generate a unique QR code on your screen. Scan this code with your smartphone to bridge your mobile device with the project. This turns your phone into a remote sensor, allowing you to capture and upload images directly to the cloud.

5. Labeling Data for Supervised Learning

Before capturing an image, it is vital to assign a Label (e.g., "cat," "dog," or "car"). Accurate labeling at the source ensures the machine learning model can correctly categorize the features of each image during the training process.

6. Dataset Partitioning (Train/Test Split)

After collecting the images, the dataset must be divided to evaluate model performance fairly. In this task, a total of 20 images were collected and split using an 80:20 ratio:

- Training Set: 16 images (used to teach the model).
- Testing Set: 4 images (used to validate accuracy).

The screenshot shows the Edge Impulse web interface. On the left is a sidebar with various project management and development tools like Dashboard, Devices, Data acquisition, Experiments, EON Tuner, Impulse design, Create impulse, Live classification, Model testing, Deployment, and an Upgrade Plan section. The main area is titled 'Dataset' and shows a summary: 'sushma0208 / sushma0208-project-1 PERSONAL' and 'Target: Cortex-M4F 80MHz'. It displays '20 items' and a 'TRAIN / TEST SPLIT' of '80% / 20%'. Below this is a table titled 'Dataset' with columns for 'SAMPLE NAME', 'LABEL', and 'ADDED'. The table lists eight entries under 'Training' and two under 'Test'. At the bottom right of the main area is a red button labeled 'Click on a sample to load...'.

SAMPLE NAME	LABEL	ADDED
Chair.6ebc1csj	Chair	Yesterday, 17...
Chair.6ebbvr1k	Chair	Yesterday, 17...
Chair.6ebbvnnfi	Chair	Yesterday, 17...
Keyboard.6ebbrkfc	Keyboard	Yesterday, 17...
Keyboard.6ebbrej9	Keyboard	Yesterday, 17...
Keyboard.6ebbrach	Keyboard	Yesterday, 17...
Desktop.6ebbpkjip	Desktop	Yesterday, 17...
Dataset.6ebbh2o4	Desktop	Yesterday, 17...

EDGE IMPULSE

sushma0208 / sushma0208-project-1 PERSONAL Target: Cortex-M4F 80MHz 5

Dashboard Devices Data acquisition Experiments EON Tuner Impulse design Create impulse Live classification Model testing Deployment

Upgrade Plan
Get access to higher job limits and more collaborators.

View plans

Dataset
DATA COLLECTED 20 items

Training 16 items **Test** 4 items

Labels in your dataset

Label	Count	Ratio
Chair	8	80% / 20% (4 / 1)
Chair.Gel	4	80% / 20% (4 / 1)
Chair.Gel	4	80% / 20% (4 / 1)
Keyboard	4	80% / 20% (4 / 1)
Keyboard	4	80% / 20% (4 / 1)
Keyboard	4	80% / 20% (4 / 1)
Mouse	4	80% / 20% (4 / 1)

SUGGESTED TRAIN / TEST SPLIT 80% / 20%

Dismiss

g your dataset.

Desktop.6ebbpkjp Desktop Yesterday, 17:2...

Dataset.6ebbh2o4 Desktop Yesterday, 17:1...

Resume tutorial

The screenshot shows the Edge Impulse web interface. A modal window titled "Dataset train / test split ratio" is open. It contains a message about training and testing data, a suggestion for an 80/20 split, and a list of labels in the dataset with their current split ratios. The labels include Chair, Chair.Gel, Keyboard, and Mouse, each with an 80% / 20% (4 / 1) split. A "SUGGESTED TRAIN / TEST SPLIT" bar is set at 80%. A "Dismiss" button is at the bottom right of the modal. The main interface shows a table of training and test samples with columns for Sample Name, Label, and Added date. A large image of a black computer mouse is displayed on the right, labeled "Mouse.6ebb8k86". A "Metadata" section below it says "No metadata." A "Resume tutorial" button is at the bottom right.

EDGE IMPULSE

sushma0208 / sushma0208-project-1 PERSONAL Target: Cortex-M4F 80MHz 5

Dashboard Devices Data acquisition Experiments EON Tuner Impulse design Create impulse Live classification Model testing Deployment

Upgrade Plan
Get access to higher job limits and more collaborators.

View plans

Dataset
DATA COLLECTED 20 items

Training 16 items **Test** 4 items

SAMPLE NAME	LABEL	ADDED
Chair.6ebc1csj	Chair	Yesterday, 17:...
Chair.6ebbvrlk	Chair	Yesterday, 17:...
Chair.6ebbvnlfi	Chair	Yesterday, 17:...
Keyboard.6ebbrkfc	Keyboard	Yesterday, 17:...
Keyboard.6ebbrej9	Keyboard	Yesterday, 17:...
Keyboard.6ebbrach	Keyboard	Yesterday, 17:...
Desktop.6ebbpkjp	Desktop	Yesterday, 17:...
Dataset.6ebbh2o4	Desktop	Yesterday, 17:...
Dataset.6ebbg929	Desktop	Yesterday, 17:...
Mouse.6ebbc286	Mouse	Yesterday, 17:...
Mouse.6ebb8k86	Mouse	Yesterday, 17:...
Mouse.6ebbaab	Mouse	Yesterday, 17:...

RAW DATA
Mouse.6ebb8k86

Metadata

No metadata.

Resume tutorial